

Ethan R. Morales

Email: ethan.morales@example.com • Phone: (415) 555-7823

Location: San Francisco, CA • LinkedIn: linkedin.com/in/ethan-r-morales • GitHub: github.com/ethanrmprompt

Professional Summary

Mid-level Prompt Engineer with 4+ years of experience designing, optimizing, and deploying large language model (LLM) applications in product environments. Proven track record delivering production-ready prompt templates, retrieval-augmented generation (RAG) systems, and cost/latency optimizations for customer-facing assistants. Strengths include prompt design and testing, prompt orchestration with LangChain/LlamaIndex, embeddings and vector search (FAISS, Pinecone), and robust evaluation frameworks (A/B tests, human-in-the-loop). Adept at partnering with product and ML engineering teams to translate requirements into measurable outcomes (accuracy, latency, API spend).

Work Experience

Prompt Engineer — Lumina Labs (SaaS AI)

San Francisco, CA • May 2022 – Present

- Led design and rollout of a modular prompt template library used across three product lines (customer support, onboarding assistant, knowledge search), increasing reuse and consistency while reducing time-to-deploy for new assistants from 4 weeks to 1 week.
- Built a RAG pipeline (vector store: Pinecone; orchestration: LangChain) for product docs that improved answer relevance (Precision@5) from 56% to 83% in internal evaluations and reduced fallback-to-human rate by 42%.
- Implemented prompt compression and token packing strategies combined with conditional prompting to reduce average API token usage per session by 38%, saving ~\$120K/year on LLM inference costs at current volume.
- Developed a prompt evaluation suite (automated unit tests + human eval interface) enabling continuous monitoring; caught and fixed 12 prompt regressions during first 6 months of use.
- Collaborated with ML engineers to design instruction-tuning data pipelines; created 18k labeled prompt-response pairs for domain alignment, resulting in a 14% lift in helpfulness metric on internal KPI.
- Mentored two junior engineers on prompt best practices, model selection, and retrieval techniques; improved team throughput for prompt experiments by 30%.

AI/Prompt Specialist — Flowlytics (Startup — Analytics + Conversational AI)

Remote • Jan 2020 – Apr 2022

- Architected a context-aware conversational assistant for B2B analytics customers using GPT-3.5 and on-prem FAISS-based retrieval, achieving a 4.6/5 average user satisfaction score over 6 months.
- Designed A/B experiments for alternative system and user prompt templates; selected winning variants that increased task completion rate by 21% and reduced average response time by 18%.
- Automated prompt versioning and rollout using Git-based workflows and PromptLayer for metrics; shortened safe-deployment cycle from days to hours.
- Collaborated with data team to label edge-case prompts and build a small instruction-tuning dataset (3.2k samples) that reduced hallucinations in domain-specific queries by 26%.
- Documented prompt playbooks and created onboarding materials that reduced ramp time for new engineers from 3 weeks to 10 days.

Data Analyst / Junior ML Engineer — Meridian Insights

Austin, TX • Jun 2018 – Dec 2019

- Built data pipelines and exploratory analyses that informed conversational UX prioritization; implemented monitoring dashboards to track usage and common failure modes for early virtual assistant prototypes.
- Implemented lightweight intent classification and slot-filling models (scikit-learn) to reduce trivial error cases before full LLM integration, contributing to a 12% reduction in support ticket volume.

Education

Skills

- Prompt Engineering & LLMs: prompt templates, system/user role engineering, chain-of-thought control, instruction tuning, prompt compression, few-shot and zero-shot design
- LLMs & Platforms: OpenAI (GPT-4/4o, GPT-3.5), Anthropic Claude, Meta Llama 2/3, Mistral, Hugging Face Transformers
- Retrieval & Embeddings: RAG architectures, embeddings (OpenAI/Hugging Face), FAISS, Pinecone, Milvus, Weaviate, LlamaIndex, LangChain
- ML Tools & Frameworks: PyTorch, Transformers, Hugging Face, scikit-learn, FastAPI
- Infrastructure & DevOps: Docker, Kubernetes, AWS (S3, Lambda, SageMaker), GCP, CI/CD, API design, Git
- Evaluation & Measurement: A/B testing, human-in-the-loop evaluation, precision/recall/Precision@k, prompt performance monitoring, cost/latency tradeoffs
- Soft Skills: cross-functional collaboration, technical writing, mentorship, product-focused experimentation

Relevant Projects

- Customer Support RAG Assistant (Lumina Labs)
- Built end-to-end RAG pipeline for multi-product knowledge base: ingest, chunking, embedding, vector index (Pinecone) and LangChain orchestration.
- Achieved Precision@5 = 83% (up from 56%) and reduced human escalation by 42%. Implemented semantic reranking and metadata-aware filtering.
- Prompt Template Library & Orchestration
- Created a versioned prompt template repository with automated tests and rollout controls. Templates cover instruction tuning, safety filters, persona/context injection, and cost-optimized fallbacks.
- Reduced prompt duplication by 70% and cut new assistant build time by 75%.
- Token Optimization & Cost Reduction Initiative
- Introduced adaptive context windowing and context summarization to limit tokens fed to the model; combined with dynamic temperature and model selection, reduced monthly inference spend by 38% while preserving response quality.
- Conversational Sales Assistant (Side Project)
- Designed a sales assistant using embeddings + GPT-4 for contextual responses, integrated with CRM (HubSpot) via API.
- Increased demo-booking conversion rate in pilot by 16% and average session length by 22%.

Certifications

- DeepLearning.AI – Generative AI with Large Language Models (Coursera) – 2023 (Certificate)
- LangChain – LangChain Fundamentals (course certificate) – 2023
- Pinecone – Vector Databases Fundamentals (training certificate) – 2022
- AWS Certified Cloud Practitioner – 2021

Selected Achievements & Metrics (At-a-glance)

- Prompt library adoption across product org: +3 product lines in 6 months
- RAG answer relevance (Precision@5): +27 percentage points (56% → 83%)
- API token cost reduction: 38% savings (~\$120K/year at scale)
- Human escalation reduction for support flows: 42%
- Team prompt experiment throughput improvement: 30%

References available upon request. Sample notebooks, prompt playbooks, and project demos available on GitHub (github.com/ethanrmprompt).